

$\lambda^\alpha \mathcal{M}^\alpha$  bytecode reference

| Mnemonic                                                                                                                                                                                                           | Encoding             |
|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------|
| BINOP +                                                                                                                                                                                                            | 01                   |
| Adds two integers, with wraparound.                                                                                                                                                                                | -2, +1               |
| BINOP -                                                                                                                                                                                                            | 02                   |
| Subtracts two integers, with wraparound.                                                                                                                                                                           | -2, +1               |
| BINOP *                                                                                                                                                                                                            | 03                   |
| Multiplies two integers, with wraparound.                                                                                                                                                                          | -2, +1               |
| BINOP /                                                                                                                                                                                                            | 04                   |
| Divides two integers, with wraparound. The result is rounded towards zero. The quotient is negative if exactly one operand is negative.                                                                            | -2, +1               |
| Raises an error if the divisor is 0.                                                                                                                                                                               |                      |
| BINOP %                                                                                                                                                                                                            | 05                   |
| Computes an integer remainder, with wraparound. The operation satisfies $(a / b) * b + (a \% b) = a$ . The remainder is negative if the first operand is negative.                                                 | -2, +1               |
| Raises an error if the divisor is 0.                                                                                                                                                                               |                      |
| BINOP <                                                                                                                                                                                                            | 06                   |
| Tests if the left operand is less than the right operand.                                                                                                                                                          | -2, +1               |
| BINOP <=                                                                                                                                                                                                           | 07                   |
| Tests if the left operand is less than or equal to the right operand.                                                                                                                                              | -2, +1               |
| BINOP >                                                                                                                                                                                                            | 08                   |
| Tests if the left operand is greater than the right operand.                                                                                                                                                       | -2, +1               |
| BINOP >=                                                                                                                                                                                                           | 09                   |
| Tests if the left operand is greater than or equal to the right operand.                                                                                                                                           | -2, +1               |
| BINOP ==                                                                                                                                                                                                           | 0a                   |
| Tests if the left operand is equal to the right operand. One of the operands must be an integer. Integers are never equal to values of other types.                                                                | -2, +1               |
| BINOP !=                                                                                                                                                                                                           | 0b                   |
| Tests if the left operand is not equal to the right operand. Unlike ==, both operands must be integers.                                                                                                            | -2, +1               |
| BINOP &&                                                                                                                                                                                                           | 0c                   |
| Tests if both integer operands are non-zero.                                                                                                                                                                       | -2, +1               |
| BINOP !!                                                                                                                                                                                                           | 0d                   |
| Tests if either of the operands is non-zero.                                                                                                                                                                       | -2, +1               |
| CONST k                                                                                                                                                                                                            | 10 [k: i32]          |
| Pushes a constant immediate k onto the stack as an integer value.                                                                                                                                                  | -0, +1               |
| STRING s                                                                                                                                                                                                           | 11 [s: i32]          |
| Pushes a string starting at offset s in the string table onto the stack.                                                                                                                                           | -0, +1               |
| SEXP s n                                                                                                                                                                                                           | 12 [s: i32] [n: i32] |
| Constructs an S-expression with n members. A string starting at offset s in the string table is used as the tag.                                                                                                   | -n, +1               |
| STI                                                                                                                                                                                                                | 13                   |
| Performs an indirect store to a variable. The first operand must be a reference to the variable. The second operand is assigned to the variable.                                                                   | -2, +1               |
| Pushes the second operand back onto the stack (for chained assignments).                                                                                                                                           |                      |
| Note: this instruction is never emitted by the Lama compiler.                                                                                                                                                      |                      |
| STA                                                                                                                                                                                                                | 14                   |
| Performs an indirect store to a variable or an aggregate. The operation is overloaded; its behavior depends on the second-to-top value on the stack, which must be either a reference to a variable or an integer: |                      |
| • If its type is a reference to a variable, this operation is equivalent to STI. In particular, it pops 2 operands and pushes 1.                                                                                   |                      |
| Note: the bytecode emitted by the Lama compiler never triggers this case as it does not contain LDA instructions.                                                                                                  |                      |

| Mnemonic                                                                                                                                                                                                                                                                                 | Encoding    |
|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------|
|                                                                                                                                                                                                                                                                                          |             |
| • If its type is an integer, this operation pops 3 operands and pushes 1. The first operand must be an aggregate: an S-expression, an array, or a string. The second operand (the integer) is an index into the aggregate. The third operand is assigned to the aggregate at this index. |             |
| The index must fall within the range from 0 (inclusive) to l (exclusive), where l is the length of the aggregate. Raises an error if the index is outside the bounds.                                                                                                                    |             |
| Pushes the third operand back onto the stack (for chained assignments).                                                                                                                                                                                                                  |             |
| JMP l                                                                                                                                                                                                                                                                                    | 15 [l: i32] |
| Sets the instruction counter to l.                                                                                                                                                                                                                                                       | -0, +0      |
| END                                                                                                                                                                                                                                                                                      | 16          |
| Marks the end of the procedure definition. When executed, returns the top value to the caller of this procedure.                                                                                                                                                                         | -1, +1      |
| RET                                                                                                                                                                                                                                                                                      | 17          |
| Returns the top value to the caller of this procedure.                                                                                                                                                                                                                                   | -1, +1      |
| DROP                                                                                                                                                                                                                                                                                     | 18          |
| Removes the top value from the stack.                                                                                                                                                                                                                                                    | -1, +0      |
| DUP                                                                                                                                                                                                                                                                                      | 19          |
| Duplicates the top value of the stack.                                                                                                                                                                                                                                                   | -1, +2      |
| SWAP                                                                                                                                                                                                                                                                                     | 1a          |
| Swaps the top two values on the stack.                                                                                                                                                                                                                                                   | -2, +2      |
| ELEM                                                                                                                                                                                                                                                                                     | 1b          |
| Looks up an element of an aggregate by its index. The first operand must be the aggregate: an S-expression, an array, or a string. The second operand must be an integer, taken as an index into the aggregate.                                                                          | -2, +1      |
| The index must fall within the range from 0 (inclusive) to l (exclusive), where l is the length of the aggregate. Raises an error if the index is outside the bounds.                                                                                                                    |             |
| LD G(m)                                                                                                                                                                                                                                                                                  | 20 [m: i32] |
| Pushes the mth global onto the stack.                                                                                                                                                                                                                                                    | -0, +1      |
| LD L(m)                                                                                                                                                                                                                                                                                  | 21 [m: i32] |
| Pushes the mth local onto the stack.                                                                                                                                                                                                                                                     | -0, +1      |
| LD A(m)                                                                                                                                                                                                                                                                                  | 22 [m: i32] |
| Pushes the mth function argument onto the stack.                                                                                                                                                                                                                                         | -0, +1      |
| LD C(m)                                                                                                                                                                                                                                                                                  | 23 [m: i32] |
| Pushes the mth variable captured by this closure onto the stack.                                                                                                                                                                                                                         | -0, +1      |
| LDA G(m)                                                                                                                                                                                                                                                                                 | 30 [m: i32] |
| Pushes a reference to the mth global onto the stack.                                                                                                                                                                                                                                     | -0, +1      |
| Note: this instruction is never emitted by the Lama compiler.                                                                                                                                                                                                                            |             |
| LDA L(m)                                                                                                                                                                                                                                                                                 | 31 [m: i32] |
| Pushes a reference to the mth local onto the stack.                                                                                                                                                                                                                                      | -0, +1      |
| Note: this instruction is never emitted by the Lama compiler.                                                                                                                                                                                                                            |             |
| LDA A(m)                                                                                                                                                                                                                                                                                 | 32 [m: i32] |
| Pushes a reference to the mth function argument onto the stack.                                                                                                                                                                                                                          | -0, +1      |
| Note: this instruction is never emitted by the Lama compiler.                                                                                                                                                                                                                            |             |
| LDA C(m)                                                                                                                                                                                                                                                                                 | 33 [m: i32] |
| Pushes a reference to the mth variable captured by this closure onto the stack.                                                                                                                                                                                                          | -0, +1      |
| Note: this instruction is never emitted by the Lama compiler.                                                                                                                                                                                                                            |             |
| ST G(m)                                                                                                                                                                                                                                                                                  | 40 [m: i32] |
| Stores a value in the mth global. Pushes the value back onto the stack (for chained assignments).                                                                                                                                                                                        | -1, +1      |
| ST L(m)                                                                                                                                                                                                                                                                                  | 41 [m: i32] |
| Stores a value in the mth local. Pushes the value back onto the stack (for chained assignments).                                                                                                                                                                                         | -1, +1      |
| ST A(m)                                                                                                                                                                                                                                                                                  | 42 [m: i32] |
| Stores a value in the mth function argument. Pushes the value back onto the stack (for chained assignments).                                                                                                                                                                             | -1, +1      |

| Mnemonic                                                                                                                                                                    | Encoding                                                              |
|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------|
| ST C(m)                                                                                                                                                                     | 43 [m: i32]                                                           |
| Stores a value in the mth variable captured by this closure. Pushes the value back onto the stack (for chained assignments).                                                | -1, +1                                                                |
| CJMPz l                                                                                                                                                                     | 50 [l: i32]                                                           |
| Sets the instruction pointer to l if the operand is zero. Otherwise, moves to the next instruction.                                                                         | -1, +0                                                                |
| CJMPnz l                                                                                                                                                                    | 51 [l: i32]                                                           |
| Sets the instruction pointer to l if the operand is non-zero. Otherwise, moves to the next instruction.                                                                     | -1, +0                                                                |
| BEGIN a n                                                                                                                                                                   | 52 [a: i32] [n: i32]                                                  |
| Marks the start of a procedure definition with a arguments and n locals. When executed, initializes locals to an empty value.                                               | -0, +0                                                                |
| Unlike CBEGIN, the procedure cannot use captured variables.                                                                                                                 |                                                                       |
| CBEGIN a n                                                                                                                                                                  | 53 [a: i32] [n: i32]                                                  |
| Marks the start of a closure definition with a arguments and n locals. When executed, initializes locals to an empty value.                                                 | -0, +0                                                                |
| Unlike BEGIN, the defined closure may use captured variables.                                                                                                               |                                                                       |
| CLOSURE l n V <sub>1</sub> (m <sub>1</sub> ) ... V <sub>n</sub> (m <sub>n</sub> )                                                                                           | 54 [l: i32] [n: i32] [[V <sub>i</sub> (m <sub>i</sub> ): varspec]; n] |
| varspec immediates are encoded as follows:                                                                                                                                  |                                                                       |
| • G(m): 00 [m: i32]                                                                                                                                                         |                                                                       |
| • L(m): 01 [m: i32]                                                                                                                                                         |                                                                       |
| • A(m): 02 [m: i32]                                                                                                                                                         |                                                                       |
| • C(m): 03 [m: i32]                                                                                                                                                         |                                                                       |
| Pushes a new closure with n captured variables onto the stack. The bytecode for the closure begins at l (given as an offset from the start of the bytecode).                | -0, +1                                                                |
| The instruction has a variable-length encoding: the description of each captured variable is specified as a 5-byte immediate.                                               |                                                                       |
| CALLC n                                                                                                                                                                     | 55                                                                    |
| Calls a closure with n arguments. The first operand must be the closure, followed by the arguments. Pushes the returned value onto the stack.                               | -(n+1), +1                                                            |
| CALL l n                                                                                                                                                                    | 56 [l: i32] [n: i32]                                                  |
| Calls a function with n arguments. The bytecode for the function begins at l (given as an offset from the start of the bytecode). Pushes the returned value onto the stack. | -n, +1                                                                |
| l must not refer to a closure definition (declared with CBEGIN) because this instruction does not capture any variables.                                                    |                                                                       |
| TAG s n                                                                                                                                                                     | 57 [s: i32] [n: i32]                                                  |
| Tests whether the operand is an S-expression with a specific tag (the string starting at offset s in the string table) and number of elements (n).                          | -1, +1                                                                |
| If the operand is not an S-expression, pushes 0.                                                                                                                            |                                                                       |
| ARRAY n                                                                                                                                                                     | 58 [n: i32]                                                           |
| Tests whether the operand is an array of elements.                                                                                                                          | -1, +1                                                                |
| FAIL ln col                                                                                                                                                                 | 59 [ln: i32] [col: i32]                                               |
| Raises an error, reporting a match failure at line ln, column col (both 1-based). The operand is the value being matched.                                                   | -1, +1                                                                |
| LINE ln                                                                                                                                                                     | 5a [ln: i32]                                                          |
| Marks the following bytecode as corresponding to line ln in the source text. Only used for diagnostics.                                                                     | -0, +0                                                                |
| PATT =str                                                                                                                                                                   | 60                                                                    |
| Tests whether the two operands are both strings and store the same bytes.                                                                                                   | -2, +1                                                                |
| PATT #string                                                                                                                                                                | 61                                                                    |
| Tests whether the operand is a string.                                                                                                                                      | -1, +1                                                                |
| PATT #array                                                                                                                                                                 | 62                                                                    |
| Tests whether the operand is an array.                                                                                                                                      | -1, +1                                                                |
| PATT #sexp                                                                                                                                                                  | 63                                                                    |
| Tests whether the operand is an S-expression.                                                                                                                               | -1, +1                                                                |

| Mnemonic                                                                                                                                                                                                                                                                                    | Encoding    |
|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------|
| PATT #ref                                                                                                                                                                                                                                                                                   | 64          |
| Tests whether the operand has a boxed representation (passed by reference).                                                                                                                                                                                                                 | -1, +1      |
| PATT #val                                                                                                                                                                                                                                                                                   | 65          |
| Tests whether the operand has an unboxed representation (passed by value).                                                                                                                                                                                                                  | -1, +1      |
| PATT #fun                                                                                                                                                                                                                                                                                   | 66          |
| Tests whether the operand is a closure.                                                                                                                                                                                                                                                     | -1, +1      |
| CALL Lread                                                                                                                                                                                                                                                                                  | 70          |
| Calls the built-in function read. The function returns the next program input. If the program input is exhausted, raises an error.                                                                                                                                                          | -0, +1      |
| Consecutive calls to read returns consecutive inputs.                                                                                                                                                                                                                                       |             |
| CALL Lwrite                                                                                                                                                                                                                                                                                 | 71          |
| Calls the built-in function write. The operand must be an integer. The function adds the operand to the program output. Returns an empty value.                                                                                                                                             | -1, +1      |
| CALL Llength                                                                                                                                                                                                                                                                                | 72          |
| Calls the built-in function length. The operand must be an aggregate: an S-expression, an array, or a string. The function returns the length of the aggregate as an integer.                                                                                                               | -1, +1      |
| CALL Lstring                                                                                                                                                                                                                                                                                | 73          |
| Calls the built-in function string. The operand must be an integer, a string, an array, or an S-expression. If the operand is an array or an S-expression, the type requirements apply transitively to the operand's elements. The function returns a string representation of the operand. | -1, +1      |
| CALL Barray n                                                                                                                                                                                                                                                                               | 74 [n: i32] |
| Calls the built-in function .array. The function creates an array composed of the n operands and returns it.                                                                                                                                                                                | -n, +1      |

## Notation

- Literal bytes in the encoding are written in hexadecimal. Integer immediates are encoded as signed numbers in native endianness.
- The number in red tells how many operands the operation pops off the stack. The number in green indicates how many values it then pushes onto the stack.

## Notes

- Arithmetic is performed modulo  $2^{31}$  on 32-bit platforms and  $2^{63}$  on 64-bit platforms. All operations are signed.
- Boolean values (resulting from comparisons) are represented as integers: 1 if true, 0 if false. For logical operations, a non-zero integer value is true.
- Operands are ordered from the lowest up; the rightmost operand is on the top.
- Operations perform type-checking dynamically, raising an error if an operand has an unexpected type.
- Jump targets are byte offsets from the start of the bytecode. In other words, all jumps are absolute.
- A closure is a procedure that captures variables in outer scopes. Variables are captured by-value, not by-reference. A closure with no captured values can be called as a regular procedure (see CALL).
- This reference assumes that no values introduced by a procedure (other than a value to be returned) remain on the stack when RET or END is executed – that is, the stack height at exit points is larger than at the entry by exactly one element.
- Integers are always stored unboxed and passed by value. All other types are always boxed and passed by reference.